

Introduction to Data Science

Classification and nonlinear models

Zhen Zhang

Southern University of Science and Technology

Outlines

Introduction

Logistic Regression

k-Nearest Neighbor

Decision Trees

Naive Bayes

Linear Discriminant Analysis

Support Vector Machine

Model Assessment

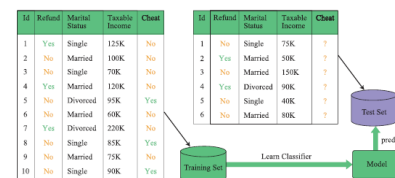
References

Why We Need Classification

- Knowing the classes of the data, we could easily manage the data and react to the possible outcomes
- Predict whether users would default in the future based on their basic information and historical transaction records
- Predict whether a tumor is benign or malignant based on their physical and geometrical features
- Predict the users' interests in the new products based on their historical purchasing records and behavioral preferences
- Separate spams and advertisements from emails

What is Classification

- Supervised learning: predict label y from features \mathbf{x}
- Training stage: Given a data set $D = \{(\mathbf{x}, y)\}$, including both features and labels, split $D = D_{\text{train}} \cup D_{\text{test}}$, find a classifier (function $y = f(\mathbf{x})$) that best relates y_{train} with $\mathbf{x}_{\text{train}}$, then evaluate how close $f(\mathbf{x}_{\text{test}})$ is to y_{test}
- Predicting stage: apply the predictor to the unlabeled data \mathbf{x}_{pred} (only features) to find the proper labels $y_{\text{pred}} = f(\mathbf{x}_{\text{pred}})$



Classification Methods

- Different assumptions on f lead to different models
- Basic classification models
 - Logistic regression
 - k-nearest neighbor (kNN)
 - Decision trees
 - Naive Bayes
 - Linear discriminant analysis (LDA)
 - Support vector machines (SVM)
 - Artificial neural network (ANN)
 - ...
- Ensemble learning: Random forest and Adaboost

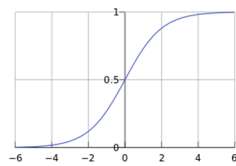
Outlines

Logistic Regression

k-Nearest Neighbor

Logistic Regression

- Not regression, but a classification method
- Connection with linear regression:
 $y = w_0 + w_1x + \epsilon$, y is binary (0 or 1); then
 $E(y|x) = P(y = 1|x) = w_0 + w_1x$; but $w_0 + w_1x$ may not be a probability
- Find a function to map it back to $[0, 1]$: Sigmoid function $g(z) = \frac{1}{1+e^{-z}}$ with $z = w_0 + w_1x_1 + \dots + w_dx_d$



- Equivalently,
 $\log \frac{P(y=1|x)}{1-P(y=1|x)} = w_0 + w_1x_1 + \dots + w_dx_d$,
 logit transform
 $\text{logit}(z) = \log \frac{z}{1-z}$

Navigation icons: back, forward, search, etc.

MLE for Logistic Regression

- The prob. distribution for two-class logistic regression model is

$$Pr(y = 1|\mathbf{X} = \mathbf{x}) = \frac{\exp(\mathbf{w}^T \mathbf{x})}{1 + \exp(\mathbf{w}^T \mathbf{x})},$$

$$Pr(y = 0|\mathbf{X} = \mathbf{x}) = \frac{1}{1 + \exp(\mathbf{w}^T \mathbf{x})}.$$

- Let $P(y = k|\mathbf{X} = \mathbf{x}) = p_k(\mathbf{x}; \mathbf{w})$, $k = 0$ or 1 . The likelihood function is defined by $L(\mathbf{w}) = \prod_{i=1}^n p_{y_i}(\mathbf{x}_i; \mathbf{w})$
- MLE estimate of \mathbf{w} : $\hat{\mathbf{w}} = \arg \max_{\mathbf{w}} L(\mathbf{w})$
- Solve $\nabla_{\mathbf{w}} \log L(\mathbf{w}) = 0$ by Newton-Raphson method

Navigation icons: back, forward, search, etc.

K-class Logistic Regression

- Extend the relative ratio of probabilities to K-class:

$$\log \frac{P(y = 1|\mathbf{X} = \mathbf{x})}{P(y = K|\mathbf{X} = \mathbf{x})} = \mathbf{w}_1^T \mathbf{x}$$

$$\log \frac{P(y = 2|\mathbf{X} = \mathbf{x})}{P(y = K|\mathbf{X} = \mathbf{x})} = \mathbf{w}_2^T \mathbf{x}$$

$$\vdots$$

$$\log \frac{P(y = K-1|\mathbf{X} = \mathbf{x})}{P(y = K|\mathbf{X} = \mathbf{x})} = \mathbf{w}_{K-1}^T \mathbf{x}$$

- Probabilistic model:

$$P(y = 1|\mathbf{X} = \mathbf{x}) = \frac{e^{\mathbf{w}_1^T \mathbf{x}}}{1 + \sum_{k=1}^{K-1} e^{\mathbf{w}_k^T \mathbf{x}}}$$

$$\vdots$$

$$P(y = K-1|\mathbf{X} = \mathbf{x}) = \frac{e^{\mathbf{w}_{K-1}^T \mathbf{x}}}{1 + \sum_{k=1}^{K-1} e^{\mathbf{w}_k^T \mathbf{x}}}$$

Navigation icons: back, forward, search, etc.

Outlines

Introduction

Logistic Regression

k-Nearest Neighbor

Decision Trees

Naive Bayes

Linear Discriminant Analysis

Support Vector Machine

Model Assessment

References

Navigation icons: back, forward, search, etc.

Introduction

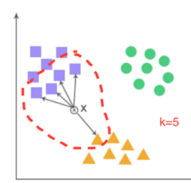
- k-nearest neighbor (kNN) is the simplest supervised learning method, especially useful when prior knowledge on the data is very limited
- Do training and test simultaneously
- When classifying a test sample x , scan the training set and find the closest k samples $D_k = \{x_1, \dots, x_k\}$ to the test sample; make vote based on the labels of the samples in D_k ; the majority vote is the label of the test sample
- Low bias, high variance
- Advantages: not sensitive to outliers, easy to implement and parallelize, good for large training set
- Drawbacks: need to tune k , take large storage, computationally intensive

Navigation icons: back, forward, search, etc.

Algorithm

- Input: training set $D_{train} = \{(x_1, y_1), \dots, (x_N, y_N)\}$, a test sample x without label y , k and distance metric $d(x, y)$
- Output: predicted label y_{pred} for x

- Compute $d(x, x_j)$ for each $(x_j, y_j) \in D_{train}$
- Sort the distances in an ascending order, choose the first k samples $(x_{(1)}, y_{(1)}), \dots, (x_{(k)}, y_{(k)})$
- Make majority vote $y_{pred} = \text{Mode}(y_{(1)}, \dots, y_{(k)})$



Navigation icons: back, forward, search, etc.

Distance Metrics

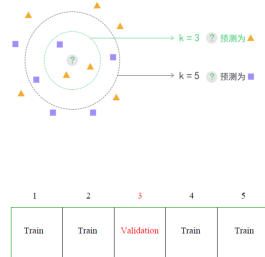
- Minkowski distance: $d_h(\mathbf{x}_1, \mathbf{x}_2) = \sqrt[h]{\sum_{i=1}^d (x_{1i} - x_{2i})^h}$; $h = 2$, Euclidean distance; $h = 1$, Manhattan distance
- Mahalanobis distance: $d(\mathbf{x}_1, \mathbf{x}_2) = \sqrt{(\mathbf{x}_1 - \mathbf{x}_2)^T \hat{\Sigma}^{-1} (\mathbf{x}_1 - \mathbf{x}_2)}$, where $\hat{\Sigma}$ is the covariance matrix of sample set; introduce correlations, could be applied to the non-scaling data
- Hamming distance: $Hamming(\mathbf{x}_1, \mathbf{x}_2) = d - \sum_{i=1}^d I(x_{1i} = x_{2i})$; used to compare two strings, e.g.,
 $Hamming('toned', 'roses') = 3$,
 $Hamming('101110', '101101') = 2$

Distance Metrics - Similarity and Divergence

- Cosine similarity: $\cos(\mathbf{x}_1, \mathbf{x}_2) = \frac{\mathbf{x}_1^T \mathbf{x}_2}{\|\mathbf{x}_1\| \|\mathbf{x}_2\|} = \frac{\sum_{i=1}^d x_{1i} x_{2i}}{\sqrt{\sum_{i=1}^d x_{1i}^2} \sqrt{\sum_{i=1}^d x_{2i}^2}}$; its range is $[-1, 1]$; the greater the cosine similarity, the more similar (closer) the two samples; insensitive to absolute value, popular in measuring user rankings; it is related to Pearson correlation coefficient
- Jaccard similarity for sets A and B : $Jaccard(A, B) = \frac{|A \cap B|}{|A \cup B|}$, used in comparing texts
- Kullback-Leibler (KL) divergence: $d_{KL}(P||Q) = \mathbb{E}_P[\log \frac{P(x)}{Q(x)}]$ measures the distance between two probability distributions P and Q ; in discrete case, $d_{KL}(p||q) = \sum_{i=1}^m p_i \log \frac{p_i}{q_i}$

Tuning k

- Different values of $k = 3$ and $k = 5$ leads to different classification results
- M -fold Cross-validation (CV) to tune k : partition the dataset into M parts ($M = 5$ or 10), let $\kappa: \{1, \dots, N\} \rightarrow \{1, \dots, M\}$ be randomized partition index map, The CV estimate of prediction error is $CV(\hat{f}, k) = \frac{1}{N} \sum_{i=1}^N L(y_i, \hat{f}^{-\kappa(i)}(x_i, k))$



Bayes Classifier (Oracle Classifier)

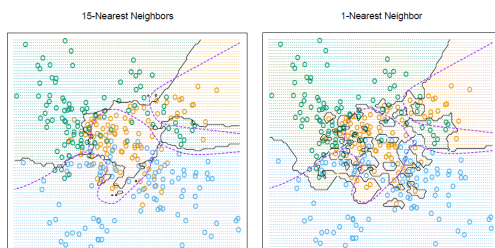
- Assume $Y \in \mathcal{Y} = \{1, 2, \dots, C\}$, the classifier $f: \mathcal{X} \rightarrow \mathcal{Y}$ is a piecewise constant function
- For 0-1 loss $L(y, f)$, the learning problem is to minimize

$$\begin{aligned} \mathcal{E}(f) &= \mathbb{E}_{P(X, Y)} L(Y, f(X)) = 1 - P(Y = f(X)) \\ &= 1 - \int_{\mathcal{X}} P(Y = f(X) | X = x) p_X(x) dx \end{aligned}$$

- Bayes rule: $f^*(x) = \arg \max_c P(Y = c | X = x)$, "the most probable label under the conditional probability on x "
- Bayes error rate: $\inf_f \mathcal{E}(f) = \mathcal{E}(f^*) = 1 - P(Y = f^*(X))$
- Bayes decision boundary: the boundary separating the K partition domains in \mathcal{X} on each of which $f^*(x) \in \mathcal{Y}$ is constant. For binary classification, it is the level set on which $P(Y = 1 | X = x) = P(Y = 0 | X = x) = 0.5$.

Decision Boundary

- The decision boundary of 15NN is smoother than that of 1NN



Analysis

- 1NN error rate is twice the Bayes error rate:
 - Bayes error = $1 - p_{c^*}(x)$ where $c^* = \arg \max_c p_c(x)$
 - Assume the samples are i.i.d., for any test sample x and small δ , there is always a training sample $z \in B(x, \delta)$ (the label of x is the same as that of z), then 1NN error is

$$\begin{aligned} \epsilon &= \sum_{c=1}^C p_c(x) (1 - p_c(x)) \xrightarrow{\delta \rightarrow 0} 1 - \sum_{c=1}^C p_c^2(x) \\ &\leq 1 - p_{c^*}^2(x) \\ &\leq 2(1 - p_{c^*}(x)) \end{aligned}$$

(Remark: In fact, $\epsilon \leq 2(1 - p_{c^*}(x)) - \frac{C}{C-1}(1 - p_{c^*}(x))^2$)

kNN Regression: Bias vs. Variance

- kNN can be used to do regression if the mode (majority vote) is replaced by mean: $\hat{f}(\mathbf{x}) = \frac{1}{k} \sum_{\mathbf{x}_{(i)} \in N_k(\mathbf{x})} y_{(i)}$
- Generalization error of kNN regression is

$$E_{train} R_{exp}(\hat{f}(\mathbf{x})) = \sigma^2 + (f(\mathbf{x}) - \frac{1}{k} \sum_{\mathbf{x}_{(i)} \in N_k(\mathbf{x})} f(\mathbf{x}_{(i)}))^2 + E_{train} \left[\frac{1}{k} \sum_{\mathbf{x}_{(i)} \in N_k(\mathbf{x})} (y_{(i)} - f(\mathbf{x}_{(i)}))^2 \right]$$

$\frac{1}{k} \sigma^2$

where we have used the fact that $E_{train} y_i = f(\mathbf{x}_i)$ and $\text{Var}(y_i) = \sigma^2$.

- For small k , overfitting, bias \searrow , variance \nearrow
- For large k , underfitting, bias \nearrow , variance \searrow

Outlines

Introduction

Logistic Regression

k-Nearest Neighbor

Decision Trees

Naive Bayes

Linear Discriminant Analysis

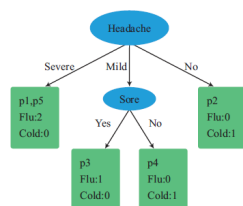
Support Vector Machine

Model Assessment

References

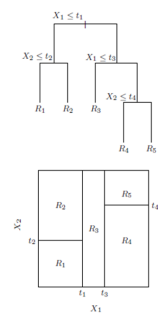
Decision Tree as Medical Diagnosis

- Diagnose whether it is flu or cold
- Rules:
 - If headache = severe, then flu
 - If headache = mild and sore = yes, then flu
 - If headache = mild and sore = no, then cold
 - If headache=no, cold



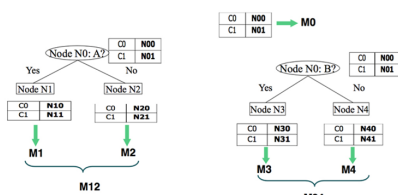
Decision Tree Algorithm

- Tree structure: internal nodes indicate features, while leaf nodes represent classes
- Start from root, choose a suitable feature x_j and its split point c_j ; at each internal node, split the node to two child nodes depending on whether $x_j \leq c_j$, until the child nodes are pure
- Equivalent to rectangular partition of the region



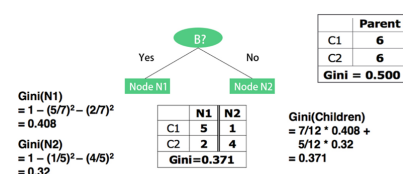
How to choose features and split points

- Impurity: choose the feature and split point so that after each split the impurity should decrease the most
- Impurity(M0)-Impurity(M12) > Impurity(M0)-Impurity(M34), choose A as split node; otherwise choose B



Impurity Measures - GINI Index

- Gini index of node t : $Gini(t) = 1 - \sum_{c=1}^C (p(c|t))^2$ where $p(c|t)$ is the proportion of class- c data in node t
- Maximum at $1 - \frac{1}{C}$, when $p(c|t) = \frac{1}{C}$
- Minimum at 0, when $p(c|t) = 1$ for some c
- Gini index of a split: $Gini_{split} = \sum_{k=1}^K \frac{n_k}{n} Gini(k)$ where n_k is the number of samples in the child node k , $n = \sum_{k=1}^K n_k$
- Choose the split so that $Gini(t) - Gini_{split}$ is maximized



Impurity Measures - Information Gain

- Entropy at t : $H(t) = -\sum_{c=1}^C p(c|t) \log_2 p(c|t)$
- Maximum at $\log_2 C$, when $p(c|t) = \frac{1}{C}$
- Minimum at 0, when $p(c|t) = 1$ for some c
- Information gain: $InfoGain_{split} = H(t) - \sum_{k=1}^K \frac{n_k}{n} H(k)$ where n_k is the number of samples in the child node k , $n = \sum_{k=1}^K n_k$
- Choose the split so that $InfoGain_{split}$ is maximized (ID3 algorithm)
- Drawback: easy to generate too many child nodes and overfit
- Introduce information gain ratio:
 $SplitINFO = -\sum_{k=1}^K \frac{n_k}{n} \log_2 \frac{n_k}{n}$, $InfoGainRatio = \frac{InfoGain_{split}}{SplitINFO}$ (C4.5 algorithm)

Algorithms

- Iterative Dichotomiser 3 (ID3): by Ross Quinlan (1986), based on Occam's Razor rule (be simple); information gain, choose feature values by enumeration
- C4.5 and C5.0: by R. Quinlan (1993), use information gain ratio instead, choose split thresholds for continuous features
- Classification and Regression Tree (CART): by Leo Breiman etc. (1984); for classification, use Gini index; for regression, use mean square error; binary split

Algorithm	Attribute Type	Impurity Measure	# Split Nodes	Target Type
ID3	Discrete	Information Gain	$k \geq 2$	Discrete
C4.5	Discrete, Continuous	Information Gain Ratio	$k \geq 2$	Discrete
C5.0	Discrete, Continuous	Information Gain Ratio	$k \geq 2$	Discrete
CART	Discrete, Continuous	GINI Index	$k = 2$	Discrete, Continuous

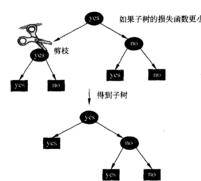
Table: Comparison of Different Decision Tree Algorithms

Tree Pruning

- Too complex tree structure easily leads to overfitting
- Prepruning: set threshold δ for impurity decrease in splitting a node; if $\Delta Impurity_{split} > \delta$, do slitting, otherwise stop
- Postpruning: based on cost function

$$Cost_{\alpha}(T) = \underbrace{\sum_{t=1}^{|T|} n_t Impurity(t)}_{\text{data fidelity}} + \alpha \underbrace{|T|}_{\text{model complexity}}$$

- Input: a complete tree T , α
- Output: postpruning tree T_{α}
 - Compute $Impurity(t)$ for $\forall t$
 - Iteratively merge child nodes bottom-up: T_A and T_B are the trees before and after merging, do merging if $Cost_{\alpha}(T_A) \geq Cost_{\alpha}(T_B)$



Comparing Three Impurity Measures

- Information gain and Gini index are more sensitive to changes in the node probabilities than the misclassification error
- Consider a two-class problem with 400 observations in each class, (400, 400); two possible splits, A: (300, 100) + (100, 300), and B: (200, 400) + (200, 0); B should be preferred
 - $Gini(A) = \frac{1}{2} Gini(A1) + \frac{1}{2} Gini(A2) = 2 \times \frac{1}{2} (2 \times \frac{3}{4} \times \frac{1}{4}) = \frac{3}{8}$,
 $Gini(B) = \frac{3}{4} Gini(A1) + \frac{1}{4} Gini(A2) = \frac{3}{4} (2 \times \frac{1}{3} \times \frac{2}{3}) = \frac{1}{3}$
 - $H(A) = 2 \times \frac{1}{2} (-\frac{3}{4} \log_2 \frac{3}{4} - \frac{1}{4} \log_2 \frac{1}{4}) = 0.81$,
 $H(B) = \frac{3}{4} (-\frac{3}{3} \log_2 \frac{1}{3} - \frac{1}{3} \log_2 \frac{2}{3}) = 0.69$
- Misclassification error at t : $Error(t) = 1 - \max_c p(c|t)$;
 $Error(A) = 2 \times \frac{1}{2} (1 - \max(\frac{3}{4}, \frac{1}{4})) = \frac{1}{4}$,
 $Error(B) = \frac{3}{4} (1 - \max(\frac{1}{3}, \frac{2}{3})) = \frac{1}{4}$
- Gini index and information gain should be used when growing the tree

ID3 Algorithm

- Input: training set $D = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$,
 $Y = \{y_1, \dots, y_n\}$, set of features $F = \{\text{column variables of } X = (\mathbf{x}_1, \dots, \mathbf{x}_n)^T\}$
- Output: decision tree T

- Create a root node
- Check Y : if all are positive, then return a single node tree T with label "+"; if all are negative, then return a single node tree T with label "-"
- Check F : if empty, then return a single node tree T with label as majority vote of Y
- For each feature in F , compute information gain, choose the feature $A \in F$ which maximizes information gain as root
- For $A = i$, let $D(i) = \{(\mathbf{x}_j, y_j) \in D | x_{jA} = i\}$:
 - If $D(i) = \emptyset$, then create a leaf node and make majority vote of D as the label
 - Else, let $D = D(i)$, go back to step 1 iteratively

Pros and Cons

- Advantages
 - Easy to interpret and visualize: widely used in finance, medical health, biology, etc.
 - Easy to deal with missing values (treat as new data type)
 - Could be extended to regression: decision tree is a rectangular partition of the domain, the predictor can be written as
 $f(x) = \sum_{m=1}^M c_m I(\mathbf{x} \in R_m)$; for regression problems
 $c_m = \bar{y}_m = \frac{1}{n_m} \sum_{i=1}^n y_i I(\mathbf{x}_i \in R_m)$ where $n_m = \sum_{i=1}^n I(\mathbf{x}_i \in R_m)$
- Drawbacks:
 - Easy to be trapped at local minimum because of greedy algorithm
 - Simple decision boundary: parallel lines to the axes

Outlines

Introduction

Logistic Regression

k-Nearest Neighbor

Decision Trees

Naive Bayes

Linear Discriminant Analysis

Support Vector Machine

Model Assessment

References

Introduction

- Based on Bayes Theorem and conditional independency assumption on features
- Widely used in text analysis, spam filtering, recommender systems, and medical diagnosis
- Bayes Theorem: let X and Y be a pair of random variables having joint probability $P(X = x, Y = y)$; by definition, the condition probability of Y given X is $P(Y|X) = \frac{P(X,Y)}{P(X)}$; then by symmetry, $P(X|Y) = \frac{P(X,Y)}{P(Y)}$; upon eliminating $P(X, Y)$

$$P(Y|X) = \frac{P(X|Y)P(Y)}{P(X)}$$

- $P(Y)$ is prior prob. distribution, $P(X|Y)$ is likelihood function, $P(X)$ is evidence, $P(Y|X)$ is posterior prob. distribution

Naive Bayes

- The core problem of machine learning is to estimate $P(Y|X)$ (or its moments $E[Y|X] = \arg \min_f E[\|Y - f(X)\|^2]$)
- Let $X = \{X_1, \dots, X_d\}$, for fixed sample $X = x$, $P(X = x)$ is independent of Y , by Bayes Theorem

$$P(Y|X = x) \propto P(X = x|Y)P(Y)$$

- Assume conditional independency of X_1, \dots, X_d given $Y = c$:

$$P(X = x|Y = c) = \prod_{i=1}^d P(X_i = x_i|Y = c)$$

- Naive Bayes model:

$$\hat{y} = \arg \max_c P(Y = c) \prod_{i=1}^d P(X_i = x_i|Y = c)$$

Maximum Likelihood Estimate (MLE)

- Estimate $P(Y = c)$ and $P(X_i = x_i|Y = c)$ from the dataset $D = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$

- MLE for $P(Y = c)$: $P(Y = c) = \frac{\sum_{i=1}^n I(y_i=c)}{n}$

- When X_i is discrete variable with range $\{v_1, \dots, v_K\}$, MLE for

$$P(X_i = v_k|Y = c) = \frac{\sum_{i=1}^n I(x_i=v_k, y_i=c)}{\sum_{i=1}^n I(y_i=c)}$$

- When X_i is continuous variable
 - Do discretization, and go back to the above formula
 - Assume X_i follows some distribution (e.g., $N(\mu, \sigma^2)$):

$$P(X_i = x|Y = c) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

Then use MLE to estimate μ and σ^2

Pros and Cons

- Where it is good
 - Spam filter: compute the posterior prob. distribution of frequently used words (convert to vector by word2vec)
 - Stable: for outliers and miss values
 - Robust: for uncorrelated features; $P(X_i|Y)$ is independent of Y and thus has no effect on posterior probability
 - May outperform far more sophisticated alternatives even if conditional independency assumption is not satisfied
- Disadvantage
 - However, when conditional independency assumption is violated, performance of Naive Bayes can be poorer
 - Depends heavily on how well the parameter estimates are

Outlines

Introduction

Logistic Regression

k-Nearest Neighbor

Decision Trees

Naive Bayes

Linear Discriminant Analysis

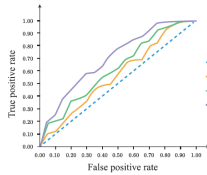
Support Vector Machine

Model Assessment

References

Receiver Operating Characteristic (ROC) and AUC

- Aim to solve class distribution imbalance problem
- Set different threshold t for continuous predicted values (probability), e.g., if $P(Y = 1|X = x_i) > t$, then $\hat{y}_i = 1$
- Compute TPR ($= \frac{TP}{TP+FN}$, or recall) vs. FPR ($= \frac{FP}{FP+TN}$) for different t and plot ROC curve
- The higher the ROC, the better the performance
- AUC: area under ROC, the larger the better, the more robust of the method for the change of t ; very good if > 0.75



Navigation icons: back, forward, search, etc.

Cohen's Kappa Coefficient (Optional)

- $\kappa = \frac{p_o - p_e}{1 - p_e} = 1 - \frac{1 - p_o}{1 - p_e}$ measures the agreement between two raters
- p_o is the accuracy (or the relative observed agreement)
- p_e is the hypothetical probability of chance agreement, $p_e = \sum_{c=1}^C \frac{n_c^{pred} n_c^{true}}{N}$, where n_c^{pred} is the number of samples predicted in class c , n_c^{true} is the true number of samples in class c , N is the total number of samples
- Eg: $p_o = \frac{20+15}{50} = 0.7$, $p_e = \frac{25}{50} \times \frac{20}{50} + \frac{25}{50} \times \frac{30}{50} = 0.5$, $\kappa = 0.4$

		Predicted Label		
		1	0	Total
True Label	1	20 TP	10 FN	30 C
	0	5 FP	15 TN	20 D
	Total	25 A	25 B	50 N

Navigation icons: back, forward, search, etc.

Other metrics

- Cohen's Kappa Coefficient $\kappa \in [-1, 1]$: as large as possible
- Multiple Classes Problem
 - ROC and AUC are not well-defined
 - Confusion matrix: $C \times C$, each entry means the number of samples in the intersection of the predicted class i and the true class j
 - Positive sample is the sample belonging to the class i , negative sample is the sample not belonging to the class i , so every sample could be positive or negative
 - Convert to multiple 0-1 classification problems
 - Precision and recall are the averages of that in the each 0-1 classification problem
 - $F1$ score is still defined as the harmonic average of precision and recall

Navigation icons: back, forward, search, etc.

Outlines

Introduction
 Logistic Regression
 k-Nearest Neighbor
 Decision Trees
 Naive Bayes
 Linear Discriminant Analysis
 Support Vector Machine
 Model Assessment
 References

Navigation icons: back, forward, search, etc.

References

- 机器学习, 周志华, 2016.
- Chapters 4, 8, 9, An Introduction to Statistical Learning with Applications in Python by Gareth James, Daniela Witten, Trevor Hastie, Robert Tibshirani and Jonathan Taylor, Springer, 2023.
- Chapters 4, 9, 12-13, 15, The Elements of Statistical Machine Learning: Data mining, Inference and Prediction by Trevor Hastie, Robert Tibshirani, and Jerome Friedman, Springer, 2009.

Navigation icons: back, forward, search, etc.