

## Introduction to Big Data Analysis Dimensionality Reduction

Zhen Zhang

Southern University of Science and Technology

## Outlines

Introduction

Principal Component Analysis


Linear Discriminant Analysis

Nonlinear Dimensionality Reduction

Feature Selection


## What is Dimensionality Reduction

- The process of reducing the number of random variables under consideration, via obtaining a set of “uncorrelated” principal variable
- By mapping from high-dimensional space to low-dimensional space
- Learning  $f: \mathcal{X} \rightarrow \mathcal{Y}$ , where  $\dim \mathcal{X} = n$  and  $\dim \mathcal{Y} = r$  with  $n > r$ .
- Including both unsupervised learning (mostly common) and supervised learning



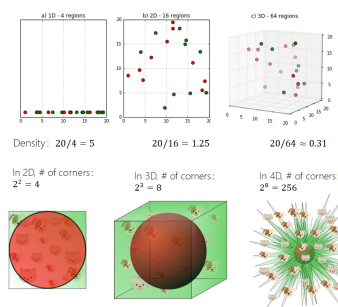
## Why Need Dimensionality Reduction ?

- Curse of dimensionality
- Eg : classify cats and dogs using features, if we want to cover 20% of the feature space, how many data do we need ?
- However, the number of samples is limited in practice




## Why Need Dimensionality Reduction ? (Cont')

- Due to the sparsity of data in high dimensions, it is easy to overfit
- Hard to train a good model to classify the corner data (getting more in high dimensions)



## Curse of Dimensionality

- The volume of hypersphere decays to zero with the increase of dimension
- The performance gets worse with the increase of dimension



## Roles of Dimensionality Reduction

- Data compression
- Denoising
- Feature extraction by mapping and feature selection (eg. Lasso)
- Reduce both spatial and time complexity, so that fewer parameters are needed and smaller computational power is required
- Data visualization

## Methods in Dimensionality Reduction

- Linear dimensionality reduction :
  - Principal component analysis (PCA)
  - Linear discriminant analysis (LDA)
  - Independent component analysis (ICA)
- Nonlinear dimensionality reduction :
  - Kernel based methods (Kernel PCA)
  - Manifold learning (ISOMAP, Locally Linear Embedding (LLE), Multidimensional scaling (MDS), t-SNE)

## Outlines

Introduction

Principal Component Analysis

Linear Discriminant Analysis

Nonlinear Dimensionality Reduction

Feature Selection

## Principal Component Analysis (PCA)


- PCA transforms a set of strongly correlated variables to another set (**typically much smaller**) of **weakly correlated** variables by using orthogonal transformation
- The new variables are called principal components
- The new set of variables are linear combinations of the original variables whose variance information is inherited as much as possible
- Unsupervised learning
- Proposed by Karl Pearson, successfully used in economics by Stone (1947) : keep 97.4% information, 17 variables about income and expenditure are finally reduced to 3 variables (F1 : total income, F2 : rate of change in total income, F3 : economic development or recession)

## Variance and Covariance Matrix

- Variance measures the variability or divergence of single variable :  $\text{Var}(X) = E(X - EX)^2$ , sample version  $S^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$ ; standard deviation :  $\text{Std}(X) = \sqrt{\text{Var}(X)}$
- For more variables,  $\text{Cov}(X, Y) = E(X - EX)(Y - EY)$ , sample version  $C = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$
- If  $X = (\mathbf{x}_1, \dots, \mathbf{x}_n)^T \in \mathbb{R}^{n \times p}$  is the sample matrix, then  $C = \frac{1}{n-1} (X - \mathbf{1}_n \bar{\mathbf{x}}^T)^T (X - \mathbf{1}_n \bar{\mathbf{x}}^T) = \frac{1}{n-1} (X - \frac{1}{n} \mathbf{1}_n \mathbf{1}_n^T X)^T (X - \frac{1}{n} \mathbf{1}_n \mathbf{1}_n^T X) = \frac{1}{n-1} X^T J X$ , where  $J = \mathbf{I}_n - \frac{1}{n} \mathbf{1}_n \mathbf{1}_n^T$  is a projection matrix with rank  $n - 1$ .

## Geometric Interpretation

- Assume a set of 2D data follows Gaussian distribution (but not limited to Gaussian distribution!), the reduction to 1D is successfully achieved by taking a direction with larger variance (larger variability of data)
- The direction in the major axis contains more information than the other direction, since smaller variance indicates the variables are almost the same



## Linear Algebra

- Let  $\{\mathbf{e}_i\}_{i=1}^p$  be the canonical basis in Euclidean space, want to find another orthonormal basis  $\{\tilde{\mathbf{e}}_i\}_{i=1}^p$  such that the random vector  $\mathbf{v} = \sum_{i=1}^p x_i \mathbf{e}_i$  can be expressed in the new basis by  $\mathbf{v} = \sum_{i=1}^p \tilde{x}_i \tilde{\mathbf{e}}_i$ , where  $\text{Var}(\tilde{x}_1) \geq \dots \geq \text{Var}(\tilde{x}_p)$  and  $\text{Cov}(\tilde{x}_i, \tilde{x}_j) \approx 0$  for  $i \neq j$
- By linear algebra, the coordinate transformation is given by the linear transformation:  $(\tilde{\mathbf{e}}_1, \dots, \tilde{\mathbf{e}}_p) = (\mathbf{e}_1, \dots, \mathbf{e}_p)W$ , where  $W \in \mathbb{R}^{p \times p}$  is an invertible matrix
- The component coefficients is transformed accordingly:  $\mathbf{x} = W\tilde{\mathbf{x}}$

◀ ▶ ↺ ↻ 🔍

## Eigendecomposition of Sample Covariance Matrix

- Assume we have  $n$  centralized samples  $\{\mathbf{x}_i\}_{i=1}^n$  with  $\frac{1}{n} \sum_{i=1}^n \mathbf{x}_i = \mathbf{0}_p$
- Then  $X^T = (\mathbf{x}_1, \dots, \mathbf{x}_n) = W(\tilde{\mathbf{x}}_1, \dots, \tilde{\mathbf{x}}_n) = W\tilde{X}^T$
- The sample covariance matrix of  $X$  is  $\text{Cov}(X) = \frac{1}{n-1} X^T X$
- The sample covariance matrix of  $\tilde{X}$  is  $\text{Cov}(\tilde{X}) = \frac{1}{n-1} \tilde{X}^T \tilde{X} = \frac{1}{n-1} W^T X^T X W = W^T \text{Cov}(X) W$
- Its diagonals are the sample versions of  $\text{Var}(\tilde{x}_1), \dots, \text{Var}(\tilde{x}_p)$ , while its off-diagonals are the covariances between  $\tilde{x}_i$  and  $\tilde{x}_j$
- Need that  $\text{Cov}(\tilde{X})$  is nearly diagonal with decreasing diagonal entries for some  $W$ .
- Equivalent to do eigendecomposition:  $\text{Cov}(X) = O \text{diag}(\lambda_1, \dots, \lambda_p) O^T$  with some orthogonal matrix  $O \in \mathbb{R}^{p \times p}$  and  $\lambda_1 \geq \dots \geq \lambda_p \geq 0$ , then let  $W = O$  completes the job

◀ ▶ ↺ ↻ 🔍

## Interpretations

- Variances in the transformed variables:  $\text{Var}(\tilde{x}_i) = \lambda_i$ , eigenvalues of  $\text{Cov}(X)$
- The new basis consists of the columns of  $W = O$ , i.e., the eigenvectors of  $\text{Cov}(X)$
- The percentage  $\frac{\lambda_i}{\sum_{j=1}^p \lambda_j}$  explains the importance of the new variable  $\tilde{x}_i$
- Given a threshold  $t$ , we can choose the number of variables  $r$  such that the total contribution to the variance of the new  $r$  variables  $\sum_{i=1}^r \frac{\lambda_i}{\sum_{j=1}^p \lambda_j}$  exceeds the threshold  $t$ . Thus these  $r$  directions  $\mathbf{w}_1, \dots, \mathbf{w}_r$  are enough to represent the original  $n$  variables
- For any random vector  $\mathbf{x} \in \mathbb{R}^p$ , the corresponding  $r$  principal components are thus  $\mathbf{w}_1^T \mathbf{x}, \dots, \mathbf{w}_r^T \mathbf{x}$

◀ ▶ ↺ ↻ 🔍

## Another Viewpoint - Best Reconstruction

- Note that the new basis  $\{\tilde{\mathbf{e}}_j\}_{j=1}^p$  is given by  $\tilde{\mathbf{e}}_j = \mathbf{w}_j$
- After the projection (if we keep the first  $r$  components), the projected point of each sample  $\mathbf{x}_i$  is  $\tilde{x}_{i,1} \mathbf{w}_1 + \dots + \tilde{x}_{i,r} \mathbf{w}_r$ , where the coordinate is given by  $\tilde{x}_{i,j} = \mathbf{w}_j^T \mathbf{x}_i$
- The reconstruction error is the sum of all squared  $L^2$  errors of all samples:

$$\begin{aligned} RE(W) &= \sum_{i=1}^n \left\| \sum_{j=1}^r \tilde{x}_{i,j} \mathbf{w}_j - \mathbf{x}_i \right\|_2^2 = \sum_{i=1}^n \left\| (W_r W_r^T - I) \mathbf{x}_i \right\|_2^2 \\ &= \sum_{i=1}^n \mathbf{x}_i^T (I - W_r W_r^T) \mathbf{x}_i = \text{Tr} \left( \sum_{i=1}^n \mathbf{x}_i \mathbf{x}_i^T (I - W_r W_r^T) \right) \\ &= \text{Tr}(X^T X (I - W_r W_r^T)) = \text{Tr}(X^T X) - \text{Tr}(W_r^T X^T X W_r) \end{aligned}$$

- Resulting in an optimization problem:  $\min_{W_r} -\text{Tr}(W_r^T X^T X W_r)$ , subject to  $W_r^T W_r = I$

◀ ▶ ↺ ↻ 🔍

## PCA Algorithm

- Given the data matrix  $X = (\mathbf{x}_1, \dots, \mathbf{x}_n)^T \in \mathbb{R}^{n \times p}$  and a threshold  $t$  (in some other cases, the number of principal components  $r$ ):
  - Centralize the data by their mean  $\bar{\mathbf{x}} = \frac{1}{n} \mathbf{1}_n^T X$ , and compute the sample covariance matrix  $C = \frac{1}{n-1} (X - \mathbf{1}_n \bar{\mathbf{x}}^T)^T (X - \mathbf{1}_n \bar{\mathbf{x}}^T)$
  - Compute the eigenvalues  $\{\lambda_i\}_{i=1}^p$  and the corresponding eigenvectors  $\{\mathbf{w}_i\}_{i=1}^p$
  - Order the eigenvalues as  $\lambda_{(1)} \geq \dots \geq \lambda_{(p)}$ , and compose an orthogonal matrix  $W$  by the eigenvectors columnwise in the same order:  $W = (\mathbf{w}_1, \dots, \mathbf{w}_p)$
  - Compute the variance contribution of the first  $r$  eigenvalues:  $\sum_{i=1}^r \frac{\lambda_i}{\sum_{j=1}^p \lambda_j}$ , find a suitable  $r$  such that this variance contribution is greater than the threshold  $t$
  - Pick the first  $r$  columns in  $W$  and form a matrix  $W_r = (\mathbf{w}_1, \dots, \mathbf{w}_r) \in \mathbb{R}^{p \times r}$
  - Output  $\tilde{X}_r = X W_r \in \mathbb{R}^{n \times r}$  as the projected data matrix, whose rows consist of data points in  $r$  dimensional subspace

◀ ▶ ↺ ↻ 🔍

## An Example

- The data: the monthly prices of three brands of vehicles (Jeep:  $x_1$ , Toyota:  $x_2$ , Benz:  $x_3$ )
- The covariance matrix is given by

$$C = \begin{pmatrix} 1 & \frac{2}{\sqrt{10}} & -\frac{2}{\sqrt{10}} \\ \frac{2}{\sqrt{10}} & 1 & -\frac{4}{5} \\ -\frac{2}{\sqrt{10}} & -\frac{4}{5} & 1 \end{pmatrix}$$

- Compute the characteristic polynomial:

$$\det(\lambda I - C) = \begin{vmatrix} \lambda - 1 & -\frac{2}{\sqrt{10}} & \frac{2}{\sqrt{10}} \\ \frac{2}{\sqrt{10}} & \lambda - 1 & \frac{4}{5} \\ \frac{2}{\sqrt{10}} & \frac{4}{5} & \lambda - 1 \end{vmatrix}$$

- Solve for the eigenvalues:  $\lambda_1 = 2.38$ ,  $\lambda_2 = 0.42$ ,  $\lambda_3 = 0.2$

◀ ▶ ↺ ↻ 🔍

## An Example (Cont')

- Plug in each eigenvalues and solve for the corresponding eigenvectors, e.g.,  $(\lambda_1 I - C)\mathbf{w}_1 = \mathbf{0}$ , or equivalently,
 
$$\begin{cases} 1.38w_{11} - \frac{2}{\sqrt{10}}w_{12} + \frac{2}{\sqrt{10}}w_{13} = 0, \\ -\frac{2}{\sqrt{10}}w_{11} + 1.38w_{12} + 0.8w_{13} = 0, \\ \frac{2}{\sqrt{10}}w_{11} + 0.8w_{12} + 1.38w_{13} = 0. \end{cases}$$
- One can find three eigenvectors as  $\mathbf{w}_1 = (0.54, 0.59, -0.59)^T$ ,  $\mathbf{w}_2 = (0.84, -0.39, 0.39)^T$ ,  $\mathbf{w}_3 = (0, 0.71, 0.71)^T$
- The three components are
 
$$\begin{aligned} \tilde{x}_1 &= \mathbf{w}_1^T \mathbf{x} = 0.54x_1 + 0.59x_2 - 0.59x_3, \\ \tilde{x}_2 &= \mathbf{w}_2^T \mathbf{x} = 0.84x_1 - 0.39x_2 + 0.39x_3, \\ \tilde{x}_3 &= \mathbf{w}_3^T \mathbf{x} = 0.71x_2 + 0.71x_3. \end{aligned}$$
- As  $\lambda_1 \gg \lambda_2, \lambda_3$ , the first principal component  $\tilde{x}_1$  reflects the change of prices in all three brands of vehicles

## Outlines

Introduction

Principal Component Analysis


Linear Discriminant Analysis

Nonlinear Dimensionality Reduction

Feature Selection

## Linear Discriminant Analysis (LDA)

- Supervised learning : based on the labels, do linear projection in order to maximize the between-class point scatter (variability) in low dimensions
- Initially proposed by R. Fisher for two-class classification (1936)
- Generalized by C. R. Rao (1948) to  $K$  classes  $\{C_1, \dots, C_K\}$



## Basic Concepts

- The number of samples in each class is  $n_k = \sum_{i:\mathbf{x}_i \in C_k} 1$ , whereas the total number of samples is  $n = \sum_{k=1}^K n_k$
- The mean of samples in class  $k$  is  $\mu_k = \frac{1}{n_k} \sum_{i:\mathbf{x}_i \in C_k} \mathbf{x}_i$ , whereas the mean of all samples is  $\mu = \sum_{k=1}^K \frac{n_k}{n} \mu_k$
- Before projection, the between-class point scatter is  $S_b = \sum_{k=1}^K \frac{n_k}{n} (\mu_k - \mu)(\mu_k - \mu)^T$ ; after projection  $W_r \in \mathbb{R}^{p \times r}$ , the between-class point scatter is  $\tilde{S}_b = W_r^T S_b W_r$
- Before projection, the within-class point scatter (variance) for each class  $C_k$  is  $S_k = \frac{1}{n_k} \sum_{i:\mathbf{x}_i \in C_k} (\mathbf{x}_i - \mu_k)(\mathbf{x}_i - \mu_k)^T$ , thus the total within-class point scatter is  $S_w = \sum_{k=1}^K \frac{n_k}{n} S_k$ ; after projection, the within-class point scatter for each class  $C_k$  is  $\tilde{S}_k = W_r^T S_k W_r$ , and the total within-class point scatter is  $\tilde{S}_w = W_r^T S_w W_r$

## Optimization Problem

- Need to find the optimal directions (columns of  $W_r$ ) such that the between-class point scatter  $\tilde{S}_b$  is maximized and within-class point scatter  $\tilde{S}_w$  is minimized, i.e.,

$$\max_{\mathbf{w}} J(\mathbf{w}) = \frac{\mathbf{w}^T \tilde{S}_b \mathbf{w}}{\mathbf{w}^T \tilde{S}_w \mathbf{w}}$$

- This is equivalent to solve

$$\max_{\mathbf{w}} J_b(\mathbf{w}) = \mathbf{w}^T \tilde{S}_b \mathbf{w}, \quad \text{subject to } \mathbf{w}^T \tilde{S}_w \mathbf{w} = 1$$

- By introducing a Lagrange multiplier  $\lambda$ , we define Lagrangian as  $L(\mathbf{w}, \lambda) = \mathbf{w}^T \tilde{S}_b \mathbf{w} - \lambda(\mathbf{w}^T \tilde{S}_w \mathbf{w} - 1)$
- The optima is obtained as the solution to the equation

$$\nabla_{\mathbf{w}} L = 2\tilde{S}_b \mathbf{w} - 2\lambda \tilde{S}_w \mathbf{w} = \mathbf{0} \Rightarrow \tilde{S}_w^{-1} \tilde{S}_b \mathbf{w} = \lambda \mathbf{w}$$

- The optimal directions are the eigenvectors of  $\tilde{S}_w^{-1} \tilde{S}_b$

## An Example

- Given two sets of data : class 1 is  $\{(4, 1)^T, (2, 4)^T, (2, 3)^T, (3, 6)^T, (4, 4)^T\}$ , and class 2 is  $\{(9, 10)^T, (6, 8)^T, (9, 3)^T, (8, 7)^T, (10, 8)^T\}$
- Class means :  $\mu_1 = (3, 3.6)^T$ ,  $\mu_2 = (8.4, 7.6)^T$ , the point scatter metrics are

$$S_1 = \begin{pmatrix} 0.8 & -0.4 \\ -0.4 & 2.6 \end{pmatrix}, \quad S_2 = \begin{pmatrix} 1.84 & -0.28 \\ -0.28 & 5.36 \end{pmatrix},$$


$$S_b = \begin{pmatrix} 7.29 & 4.86 \\ 4.86 & 3.24 \end{pmatrix}, \quad S_w = \begin{pmatrix} 1.32 & -0.34 \\ -0.34 & 4 \end{pmatrix}.$$

- The eigenvalue of  $\tilde{S}_w^{-1} \tilde{S}_b$  is solved from

$$0 = \det(\lambda I - \tilde{S}_w^{-1} \tilde{S}_b) = \begin{vmatrix} \lambda - 5.97 & -3.98 \\ -1.72 & \lambda - 1.15 \end{vmatrix} \Rightarrow \lambda = 7.11$$

## An Example (Cont')

- The optimal directions is  $\mathbf{w}^* = (0.96, 0.28)^T$
- After projection, the data become 1D :
  - Class 1 : {4.12, 3.03, 2.75, 4.55, 4.95}
  - Class 2 : {11.42, 7.98, 9.48, 9.63, 11.83}



## PCA vs. LDA

Rk: 都是线性降维

- PCA
  - Start from sample covariance matrix and find directions with maximal variances
  - Unsupervised learning, used as pre-training step, must be coupled with other learning methods
- LDA
  - Make use of labels and find projections after which the classification becomes more obvious
  - Supervised learning, can be used as classification or coupled with other learning methods

## Outlines

Introduction

Principal Component Analysis


Linear Discriminant Analysis

Nonlinear Dimensionality Reduction

Feature Selection

## Kernel PCA

- PCA works well for Gaussian distribution
- If the data do not follow Gaussian, we can find a map  $\phi: \mathbb{R}^p \rightarrow \mathbb{R}^q$  so that  $\phi(\mathbf{x})$  (almost) follows Gaussian
- We can do PCA for the transformed data  $\{\phi(\mathbf{x}_i)\}_{i=1}^n$
- Similar to nonlinear SVM, kernel trick can be used to avoid explicit computation of  $\phi$



## Covariance Matrix in Transformed Space

- Assume the transformed data are centralized :  $\mu = \frac{1}{n} \sum_{i=1}^n \phi(\mathbf{x}_i) = 0$
- Covariance Matrix  $\tilde{C} = \frac{1}{n-1} \sum_{i=1}^n \phi(\mathbf{x}_i) \phi(\mathbf{x}_i)^T$
- Do PCA for transformed data is equivalent to find the eigenvalues and eigenvectors of  $\tilde{C}$   $\tilde{C} \tilde{v} = \frac{1}{n-1} \sum_{i=1}^n \phi(\mathbf{x}_i) \phi(\mathbf{x}_i)^T \tilde{v} = \frac{1}{n-1} \sum_{i=1}^n \phi(\mathbf{x}_i) \tilde{v}^T \phi(\mathbf{x}_i)$
- Let  $\lambda$  be an eigenvalue of  $\tilde{C}$  and  $\mathbf{v} \in \mathbb{R}^q$  be the corresponding eigenvector, i.e.,  $\tilde{C} \mathbf{v} = \lambda \mathbf{v}$ .  $\lambda \tilde{v} = \frac{1}{n-1} \sum_{i=1}^n \phi(\mathbf{x}_i) \phi(\mathbf{x}_i)^T \tilde{v}$
- It can be shown that  $\mathbf{v} = \sum_{i=1}^n \alpha_i \phi(\mathbf{x}_i)$  where  $\lambda \tilde{v} = \frac{1}{n-1} \sum_{i=1}^n \phi(\mathbf{x}_i) \phi(\mathbf{x}_i)^T \tilde{v}$   $\alpha_i = \frac{1}{\lambda(n-1)} \sum_{j=1}^n \phi(\mathbf{x}_j)^T \tilde{v}$
- Furthermore,  $\alpha_i = \frac{1}{\lambda(n-1)} \sum_{j=1}^n K(\mathbf{x}_i, \mathbf{x}_j) \alpha_j$ , where  $K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$  is kernel function  $\Rightarrow \alpha_i = \frac{1}{\lambda(n-1)} \sum_{j=1}^n \phi(\mathbf{x}_j)^T \tilde{v}$
- It is sufficient to solve the eigenvalue problem :  $K \alpha = \lambda(n-1) \alpha$  where  $K = (K(\mathbf{x}_i, \mathbf{x}_j))_{i,j}$  is kernel matrix and  $\alpha = (\alpha_i)$  is the coefficient vector of  $\mathbf{v}$

## Kernel PCA Algorithm

- Choose a kernel function  $K(x, y)$  satisfying the necessary properties
- Compute the kernel matrix  $K = (K(\mathbf{x}_i, \mathbf{x}_j))_{i,j}$
- Compute the eigenvalues  $\lambda_1 \geq \dots \geq \lambda_q$  and eigenvectors  $\alpha^{(1)}, \dots, \alpha^{(q)}$  of  $K$
- For any new sample  $\mathbf{x}$ , the  $j$  component after projection is


$$z_j = \mathbf{v}_j^T \phi(\mathbf{x}) = \sum_{i=1}^n \alpha_i^{(j)} K(\mathbf{x}_i, \mathbf{x})$$

## Kernel PCA : An Example

### Wine data (from UCI repository)

13 dim. chemical measurements of for three types of wine. 178 data.  
Class labels are **NOT** used in PCA, but shown in the figures.


First two principal components:



Navigation icons: back, forward, search, etc.

## Manifolding Learning


- A manifold is a topological space that locally resembles Euclidean space near each point. It generalizes the concepts of curves and surfaces in Euclidean space.
- The dimension of a manifold is the minimal number of coordinates to represent a point on the manifold
- Some dimensionality reduction methods are based on the concept of manifold : ISOMAP, LLE, MDS, t-SNE



Navigation icons: back, forward, search, etc.

## Locally Linear Embedding (LLE)

- Reduce the number of free coordinates while keeping the local geometric structure of the data, e.g., if  $\mathbf{x}_A$  and  $\mathbf{x}_B$  are neighbor in high dimension, after the dimension reduction (transformation), they must be close to each other in low dimension
- The clustering effect should also be inherited



Navigation icons: back, forward, search, etc.


## LLE Reconstruction

- Assume each data point is locally linearly dependent of its neighbors : it can be written as the linear combination of its  $K$  nearest neighbors  $\{\mathbf{x}_{kij}\}_{j=1}^K$ , with the KNN indices  $\{k_{ij}\}_{j=1}^K$
- The weight is determined by the optimization for each  $\mathbf{x}_i$  :

$$\min_{\mathbf{w}} \|\mathbf{x}_i - \sum_{j=1}^K w_{kij} \mathbf{x}_{kij}\|_2^2$$

$$\text{subject to } \sum_{j=1}^K w_{kij} = 1, \quad w_{ij} \geq 0$$

where  $w_{ij} = 0$  if  $j \notin \{k_{ij}\}_{j=1}^K$



Navigation icons: back, forward, search, etc.

## Low Dimensional Representation

- In  $r$  ( $r < p$ ) dimensional space, find  $n$  points such that the local structure (e.g., clustering effect) is preserved

$$\min_{\mathbf{y}_1, \dots, \mathbf{y}_n} \sum_{i=1}^n \|\mathbf{y}_i - \sum_{j=1}^n w_{ij} \mathbf{y}_j\|_2^2$$

- This is equivalent to the matrix minimization problem

$$\min_{\mathbf{Y}} \text{Tr}(\mathbf{Y}^T \mathbf{M} \mathbf{Y}), \quad \text{s.t. } \mathbf{Y} \mathbf{Y}^T = \mathbf{I},$$

where  $\mathbf{Y} = (\mathbf{y}_1, \dots, \mathbf{y}_n)^T \in \mathbb{R}^{n \times r}$  and  $\mathbf{M} = (\mathbf{I} - \mathbf{W})^T (\mathbf{I} - \mathbf{W})$  with  $\mathbf{W} = (w_{ij})_{i,j=1}^n$  being the weight matrix (not necessarily symmetric)

- This is solved by eigen-decomposition : The columns of  $\mathbf{Y}$  consist of the  $r$  eigenvectors corresponding to the  $r$  smallest eigenvalues of  $\mathbf{M}$

Navigation icons: back, forward, search, etc.


## Summary of LLE

- Only one tuning parameter  $K$
- Linear algebra computation
- Only local information, no global information
- No explicit mapping as in PCA ( $\tilde{\mathbf{X}}_r = \mathbf{X} \mathbf{W}_r$ )

Navigation icons: back, forward, search, etc.

## The Motivation of ISOMAP


- The distance between two points may be different in different metrics (manifold metric vs. Euclidean metric)
- Geodesic distance could be a good metric instead of Euclidean distance
- Computation of geodesic distance, minimal path in graph



Navigation icons: back, forward, search, etc.

## ISOMAP Algorithm

- Construct KNN graph  $G = (V, E)$ :
  - For each  $\mathbf{x}_i$ , find its  $K$  nearest neighbors  $\{\mathbf{x}_j\}_{j \in N(i)}$
  - The weight of the edge  $\langle i, j \rangle$  between  $\mathbf{x}_i$  and  $\mathbf{x}_j$  is the Euclidean distance for each  $j \in N(i)$
  - Use Floyd algorithm to compute the minimal path between each pair of vertices  $(i, j)$  as the geodesic distance  $d_G(i, j)$
  - Find the low dimensional representation (e.g. by MDS):



$$\min_{\mathbf{y}_1, \dots, \mathbf{y}_n} \sum_{i \neq j} (d_G(i, j) - \|\mathbf{y}_i - \mathbf{y}_j\|)^2$$

将  $d(i, j)$  换成  $d_G(i, j)$ : 高维空间中的距离可能出现问题

Navigation icons: back, forward, search, etc.

## Floyd Algorithm (Complexity $O(n^3)$ )

### 1. Initialization :

$$d_G(i, j) = \begin{cases} d_x(i, j), & \text{if } \langle i, j \rangle \in E \\ \infty, & \text{otherwise} \end{cases}$$

- For each pair  $(i, j)$ , update the distance as follows : for each  $k = 1, \dots, n$ ,  $d_G(i, j) = \min\{d_G(i, j), d_G(i, k) + d_G(k, j)\}$
- The final output  $d_G(i, j)$  is the geodesic distance between  $i$  and  $j$

Navigation icons: back, forward, search, etc.

## Summary of ISOMAP

- Only one tuning parameter  $K$
- High computational power
- Preserve the global information
- Sensitive to noise

Navigation icons: back, forward, search, etc.

## Multidimensional Scaling (MDS)

- For data points in high dimensional space,  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^p$ , find the distance or dissimilarity matrix  $\{d_{ij}\}_{i,j=1}^n$ , e.g.,  $d_{ij} = \|\mathbf{x}_i - \mathbf{x}_j\|$
- Find  $\{\mathbf{y}_i\}_{i=1}^n \subset \mathbb{R}^r$  ( $r < p$ ), such that the distance information is preserved :

$$\min_{\mathbf{y}_1, \dots, \mathbf{y}_n} S_M(\mathbf{y}_1, \dots, \mathbf{y}_n)$$

where  $S_M(\mathbf{y}_1, \dots, \mathbf{y}_n) = \sum_{i \neq j} (d_{ij} - \|\mathbf{y}_i - \mathbf{y}_j\|)^2$  is the stress function. This is called least square or Kruskal-Shephard scaling

- Alternative objective function (Sammon mapping) :  $S_S(\mathbf{y}_1, \dots, \mathbf{y}_n) = \sum_{i \neq j} \frac{(d_{ij} - \|\mathbf{y}_i - \mathbf{y}_j\|)^2}{d_{ij}^2}$  takes care of small  $d_{ij}$
- This is nonconvex minimization

Navigation icons: back, forward, search, etc.

## t-distributed Stochastic Neighbor Embedding (t-SNE)

- Developed by Laurens van der Maaten and Geoffrey Hinton
- Effective for data visualization in 2D and 3D, applications in computer security research, music analysis, cancer research, especially for bioinformatic data
- Often display clusters in low dimensional space (may be false findings)
- With special parameter choices, approximates a simple form of spectral clustering

Navigation icons: back, forward, search, etc.

## Similarity in High Dimensional Space

- For data points in high dimensional space,  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^P$ , find the similarity of  $\mathbf{x}_i$  and  $\mathbf{x}_j$  in the form of probability  $p_{ij}$
- The similarity of data point  $\mathbf{x}_j$  to data point  $\mathbf{x}_i$  is the conditional probability,  $p_{j|i}$ , that  $\mathbf{x}_i$  would pick  $\mathbf{x}_j$  as its neighbor if neighbors were picked in proportion to their probability density under a Gaussian centered at  $\mathbf{x}_i$  :

$$p_{j|i} = \frac{\exp(-\|\mathbf{x}_i - \mathbf{x}_j\|^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-\|\mathbf{x}_i - \mathbf{x}_k\|^2 / 2\sigma_i^2)}$$


- $p_{ij} = (p_{j|i} + p_{i|j}) / 2n$ ,  $p_{ii} = 0$
- The bandwidth is adapted to the density of the data : smaller values of  $\sigma_i$  are used in denser parts of the data space

◀ ▶ ↺ ↻ 🔍

## Similarity in Low Dimensional Space

- t-SNE aims to learn a set of low dimensional data  $\mathbf{y}_1, \dots, \mathbf{y}_n \in \mathbb{R}^r$  that reflects the similarity  $p_{ij}$  as well as possible
- The similarity between the data point  $\mathbf{y}_i$  and  $\mathbf{y}_j$  follows t-distribution : (assume  $q_{ii} = 0$ )

$$q_{ij} = \frac{(1 + \|\mathbf{y}_i - \mathbf{y}_j\|^2)^{-1}}{\sum_{k \neq i} (1 + \|\mathbf{y}_k - \mathbf{y}_i\|^2)^{-1}}$$

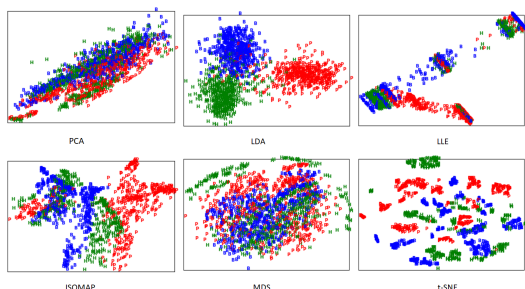


- t-distribution is heavy tailed so that large  $p_{ij}$  (dissimilar data pair) leads to even larger  $q_{ij}$  (falls apart)
- The closedness between the two similarity measures  $p_{ij}$  and  $q_{ij}$  is given by the Kullback-Leibler divergence :

$$D_{KL}(P \| Q) = \sum_{i \neq j} p_{ij} \log \frac{p_{ij}}{q_{ij}}$$

◀ ▶ ↺ ↻ 🔍

## Comparison (Optical Character Recognition)



◀ ▶ ↺ ↻ 🔍

## Outlines

Introduction

Principal Component Analysis

Linear Discriminant Analysis


Nonlinear Dimensionality Reduction

Feature Selection

◀ ▶ ↺ ↻ 🔍

## What is Feature Selection

- Low computational cost, better accuracy (avoid overfitting), and better interpretation,
- Feature engineering : feature extraction and selection. Feature extraction is according to the knowledge of the professions, usually done by expertise in the professional areas
- Three types : Filter, Wrapper, and Embedded



◀ ▶ ↺ ↻ 🔍

## Subset Selection

- Subset search :
  - Forward search (forward stepwise, forward stagewise) :  $\emptyset \Rightarrow \{x_1\} \Rightarrow \{x_1, x_4\} \Rightarrow \dots$
  - Backward search (backward stepwise) :  $\{x_1, x_2, \dots, x_p\} \Rightarrow \{x_1, x_2, \dots, x_p\} \setminus \{x_4\} \Rightarrow \dots$
  - Bidirectional search
- Evaluation metrics :
  - Distances : Euclidean, Manhattan, point scatter matrices, Kullback-Leibler divergence, etc.
  - Information : mutual information, information gain (IG), etc.
  - Correlations : Pearson correlation, Maximal information coefficients (MIC)
- Stopping rules : number of features, number of iterations, non-incremental metrics, attaining optimality, etc.
- Validation and comparison

◀ ▶ ↺ ↻ 🔍



## Three Types of Feature Selection

- Filter : filter the features by their correlations (or MIC, IG) with response variables
- Wrapper : use accuracy, precision, recall, AUC, etc.
  - Akaike Information Criteria (AIC) :  $AIC = -2\ln(L) + 2k$
  - Bayes Information Criteria (BIC) :  $BIC = -2\ln(L) + k\ln(n)$
  - Minimize AIC or BIC, where  $L$  is likelihood function,  $k$  is the number of features (parameters),  $n$  is the number of samples
- Embedded :
  - Random forest : feature importance
  - Regularization : Ridge and LASSO
  - Recursive feature elimination (RFE) : select the best (worst) feature according to the coefficients (e.g. linear regression), then do this recursively to find the feature importance

## References

- 数据分析导论，博雅大数据学院
- 周志华，机器学习，2016
- T. Hastie, R. Tibshirani, and J. Friedman, The Elements of Statistical Learning : Data mining, Inference, and Prediction, 2nd Edition, 2009
- Arthur, D., Vassilvitskii, S. "k-means++ : the advantages of careful seeding". Proceedings of the eighteenth annual ACM-SIAM symposium on Discrete algorithms, Society for Industrial and Applied Mathematics Philadelphia, PA, USA. pp. 1027 - 1035, 2007
- Lingras P, West C, Interval Set Clustering of Web Users with Rough Kmeans, Journal of Intelligent Information Systems 23(1) :5 - 16, 2004